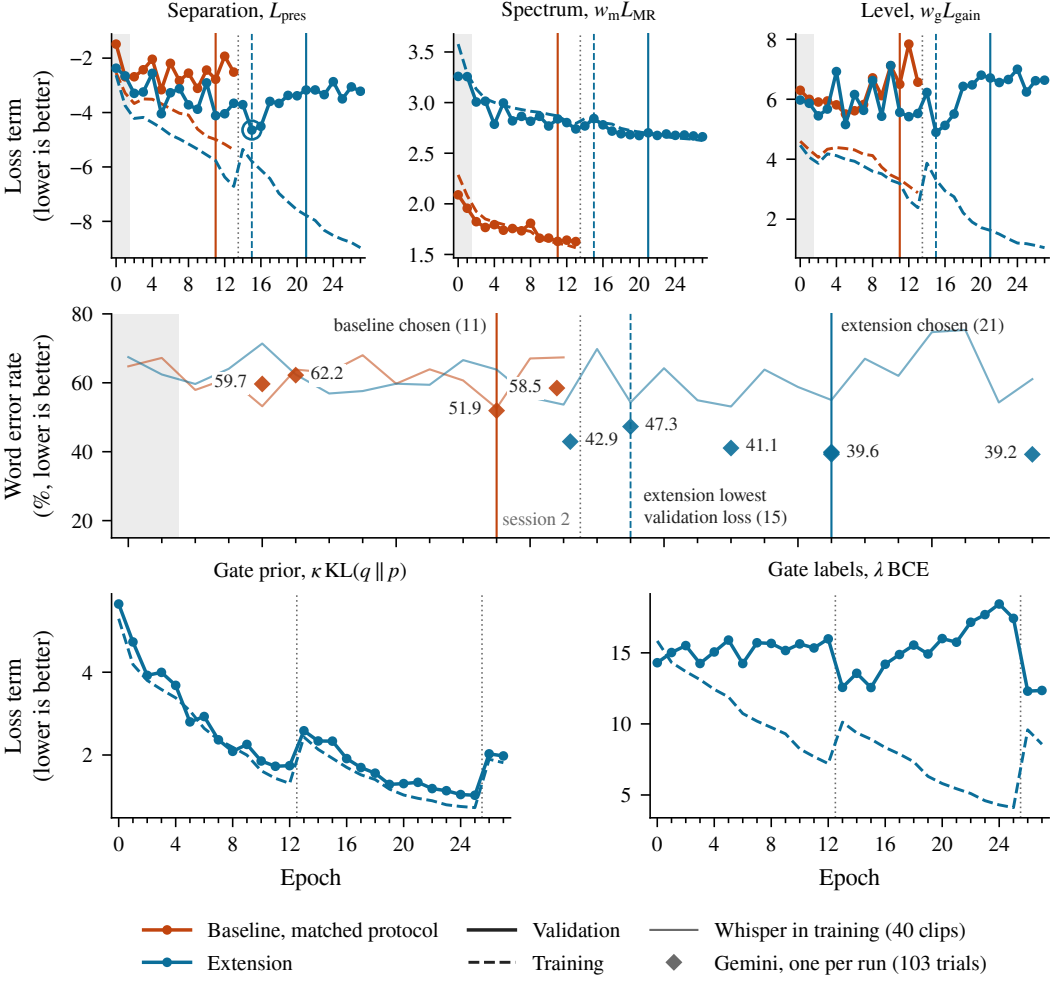
Separation, L_{pres}
Spectrum, $w_{\text{m}}L_{\text{MR}}$
Level, $w_{\text{g}}L_{\text{gain}}$
Loss term (lower is better)
Word error rate (%, lower is better)
baseline chosen (11)
extension chosen (21)
59.7
62.2
51.9
58.5
42.9
47.3
41.1
39.6
39.2
session 2
extension lowest validation loss (15)
Gate prior, $\kappa\,\mathrm{KL}(q \,\|\, p)$
Gate labels, $\lambda\,\mathrm{BCE}$
Epoch
Baseline, matched protocol
Extension
Validation
Training
Whisper in training (40 clips)
Gemini, one per run (103 trials)